



Finding Double-Unlock Bugs with Shape-and-Effect Analysis

Anders Fischer-Nielsen



Why?

- The Linux kernel is everywhere (embedded devices, routers, TVs, phones, PCs & large servers)
- Bugs in the Linux kernel can therefore have a big impact on many devices
- Shared-memory concurrency and locks are used extensively in the source code of the kernel
 - Allows parallelization of subsystems while avoiding race conditions
- Static analysis of the Control Flow Graph (CFG) of the kernel allows detecting possible resource manipulation errors



Double-Unlock Errors

- A thread holding a lock which then releases this lock more than once consecutively will result in **undefined behaviour**, according to the POSIX standard

“The results are undefined if the lock is not held by the calling thread. [...] The results are undefined if this function is called with an uninitialized thread spin lock.”

- Undefined behaviour at the kernel level possibly makes the operating system behave in unexpected ways
- A program depending on undefined behaviour might not break today, but could break in the future, since no assumptions can be made about undefined behaviour
- Detecting these errors allows developers to detect errors in their code, hopefully leading to safer programs --- provided the developers fix detected errors.



An Example

```
#include "ubifs.h"

static void orphan_delete(struct ubifs_info *c, struct ubifs_orphan *orph)
{
    if (orph->del) {
        spin_unlock(&c->orphan_lock);
        return;
    }
    if (orph->cmt) {
        spin_unlock(&c->orphan_lock);
        return;
    }
}

void ubifs_delete_orphan(struct ubifs_info *c, ino_t inum)
{
    orphan_delete(c, orph);
    spin_unlock(&c->orphan_lock);
}
```



What I Wanted to Accomplish

- *What is the performance of a double-unlock error checker prototype implementation when run on the source code of Linux kernel components?*
- Accomplished by:
 - Developing an error checker able to find these double-unlock errors
 - Analyzing the checker performance by running the implementation on confirmed double-unlock bugs found in the Linux kernel




Building Upon an Existing Implementation

- Building upon the shape-and-effect inference system described by Abal et. al. to develop a double-unlock checker relatively quickly
- Computational Tree Logic (CTL) is used by Abal et. al. to show which errors should be detected by a **double-lock** checker
- This specification can then be used to formulate a **double-unlock** checker
- CTL models program executions in time as tree-like structures where all paths can be the actual path executed at runtime


Scope

- Internal abstractions are merely used --- not modified --- in the project implementation



What We Want to Detect


$$\top \text{ EU } (\text{unlock}_p \wedge \text{EX } (\neg \text{lock}_p \text{ EU } \text{unlock}_p))$$



In other words; an unlock is performed on p without a lock on p being present before another unlock is performed on p , leading to a double-unlock

What We Want to Detect

$$\begin{aligned} & \neg \text{EU}(\text{unlock}_p \wedge \text{EX}(\neg \text{lock}_p \text{EU} \text{unlock}_p)) \\ & \text{E}[\top \cup \text{unlock}_p \wedge \text{EX} \text{E}[\neg \text{lock}_p \cup \text{unlock}_p]] \\ & \text{EF} \text{unlock}_p \wedge \text{EX} \text{E}[\neg \text{lock}_p \cup \text{unlock}_p] \end{aligned}$$



In other words; an unlock is performed on p without a lock on p being present before another unlock is performed on p , leading to a double-unlock



Implementation

- EBA framework requires implementing a **Spec** OCaml module. If all checks return an Option value, an error has been found.
- The checks found in the **Spec** signature have been implemented using existing EBA infrastructure & helper functions
- Reachability of a node matching the checker definitions is tested internally in the EBA infrastructure

```

module type Spec = sig
  (** A name to identify the checker *)
  val name : string

  (** Checker's internal state, eg. memory regions to track. *)
  type st
  (** Selects initial contexts *)
  val select : AFile.t -> Cil.fundec -> shape scheme -> AFun.t -> st L.t
  (** Flags steps of interest for triaging. *)
  val trace : st -> Effects.t -> bool
  (** Tests *)
  val testP1 : st -> step -> st option
  val testQ1 : st -> step -> st option
  val testP2 : st -> step -> st option
  (** Q2 = P2 /\ Q2-weak *)
  val testQ2_weak : st -> step -> st option

  (** Bug data *)
  type bug
  val bug_of_st : st -> bug
  val doc_of_report : fn:Cil.varinfo -> bug -> loc1:Cil.location -> loc2:Cil.location -> trace:path -> PP.doc
end

```

```

module type Spec = sig
  (** A name to identify the checker *)
  val name : string

  (** Checker's internal state, eg. memory regions to track. *)
  type st
  (** Selects initial contexts *)
  val select : AFile.t -> Cil.fundec -> shape scheme -> AFun.t -> st L.t
  (** Flags steps of interest for triaging. *)
  val trace : st -> Effects.t -> bool
  (** Tests *)
  val testP1 : st -> step -> st option
  val testQ1 : st -> step -> st option
  val testP2 : st -> step -> st option
  (** Q2 = P2 /\ Q2-weak *)
  val testQ2_weak : st -> step -> st option


  (** Bug data *)
  type bug
  val bug_of_st : st -> bug
  val doc_of_report : fn:Cil.varinfo -> bug -> loc1:Cil.location -> loc2:Cil.location -> trace:path -> PP.doc
end

```

These functions model parts of the CTL specification

Implementation

- **testP1** - Lifts a given CFG node to an Option type
- **testQ1** - Checks whether an unlock is performed without a lock in a step
- **testP2** - Checks that no lock is performed in a given step
- **testQ2_weak** - Checks whether an unlock is present in a step





Implementation

<https://github.com/lagoAbal/eba/pull/5>

```
let testQ1 st step =  
  let unlocks_and_not_locks r ef =  
    let unlocks r ef :bool = E.(mem (unlocks ~r) ef) in  
    unlocks r ef && not_locks r ef in  
  
  unlocks_and_not_locks st.reg step.offs =>?  
    let unlock_opt = find_unlock_object step in  
    let krs = Option.Infix.(  
      (unlock_opt >= Lenv.kregions_of st.fna)  
      |? Regions.empty  
    ) in  
    {st with unlock = unlock_opt; kreg = krs}
```

```
let testP2 st step =  
  let res = not_locks st.reg  
  step.offs =>? st in  
  res  
  
  let testQ2_weak st step =  
    let must_unlock = E.(mem (unlocks  
~r:st.reg) step.offs) in  
    if must_unlock  
    then Some st  
    else None
```



Implementation Caveats

- CTL formulations do not always match implementation in existing checkers
- The implementation of reachability checking in EBA is hard to reason about
- Seemingly sound checks in functions do not always give expected results
 - Due to parsing and may/must handling in EBA

More on this later...



Results

File	Status
drivers/block/drbd/drbd_main.c	Detected
fs/ubifs/orphan.c	Detected
drivers/gpu/drm/nouveau/nouveau_svm.c	Undetected
fs/btrfs/file.c	Undetected
drivers/staging/wilc1000/wilc_wlan.c	Detected
drivers/staging/kpc2000/kpc_dma/fileops.c	Detected
fs/nfs/client.c	Undetected
fs/btrfs/file.c	Undetected
drivers/media/dvb-core/dvbdev.c	Stack overflow
mm/memory_hotplug.c	Undetected
sound/soc/codecs/pcm512x.c	Detected
drivers/target/target_core_user.c	Uncaught exception
drivers/rpmsg/qcom_smd.c	Compilation error (GCC)
drivers/scsi/aacraid/commsup.c	Compilation error (GCC)
drivers/staging/rtl8188eu/os_dep/usb_intf.c	Compilation error (GCC)
block/blk-cgroup.c	Compilation error (GCC)

Technically all of these result in an EBA
ParseError without mitigations



Results

- EBA abstractions do not support all lock types found in the Linux kernel

```
if Opts.all_lock_types()
  then begin
    add_axiom tbl ax_mutex_lock;
    add_axiom tbl ax_mutex_lock_nested;
    add_axiom tbl ax_mutex_lock_interruptible_nested;
    add_axiom tbl ax_mutex_unlock;
    add_axiom tbl ax_raw_read_lock;
    add_axiom tbl ax_raw_read_unlock;
    add_axiom tbl ax__raw_read_unlock;
  end
```

- EBA parsing does not support all C constructs found in the Linux kernel
- EBA implementation of the CTL subset does not allow expressing desired checks

Where to Go From Here?

- Extending capabilities of EBA framework
 - Develop more checkers
 - Supporting lock types
 - Extend C parsing
- Reimplementing existing and future checkers as monitor state machines
 - Existing CTL implementation of EBA might limit expressiveness in checkers (e.g. `may/must`)
 - Gives greater “expressibility” of checker specifications (e.g. `if/else if`)
- Reachability of EBA can be modelled as state machines
- Current CTL can be modelled as state machines

